

Chapter 11 | Public-Key Encryption

Tuesday, November 7, 2023 11:44 AM

11.1 Public-Key Encryption—An Overview

Story:

- The introduction of public-key encryption marked a revolution in cryptography.
 - Until that time cryptographers had relied exclusively on shared secret keys to achieve private communication.
 - Public-key techniques in contrast enabled parties to communicate privately without having agreed on any secret information in advance.
 - As we have already noted it is quite amazing and counterintuitive that this is possible it means that two people on opposite sides of a room who can only communicate by shouting to each other and have no initial secret can talk in such a way that no one else in the room learns anything about what they are saying.
- In the setting of private key encryption two parties agree on a secret key that can be used by either party for both encryption and decryption
 - Public key encryption is *asymmetric* in both these respects
 - On party (the receiver) generates a pair of keys (pk, sk) called the public key and the private key resp.
 - The public key is used by a sender to encrypt a message the receiver uses the private key to decrypt the resulting ciphertext.
- Since the goal is to avoid the need for two parties to meet in advance to agree on any information how does the sender learn pk ?
[ME: Here, the emphasis is that the channel is assumed to **authenticated and public**]

At an abstract level
this can happen in two ways:

- Call the receiver Alice and
- the sender Bob

- In the first approach
when Alice learns that Bob wants to communicate with her
She can at that point generate
 pk, sk
(assuming she hasn't done so already)
and
then send pk to Bob in *the clear*.

Bob can then use pk to encrypt his message.

We emphasise that the channel between Alice and Bob may
be public
but is assumed to be **authenticated**
meaning that the adversary cannot modify the public key
sent by Alice to Bob
(and in particular cannot replace it with its own key)

See Section 12.7 for a discussion of how public keys can
be distributed over **unauthenticated channels**.

- An alternative approach
is for Alice to generate her keys (pk, sk) in advance
independently of any particular sender
(in fact, at the time of key generation
Alice need not even be aware that
Bob wants to talk to her
or even that Bob exists).
- Alice can widely disseminate her public key pk
by, say, publishing it on her webpage
putting it on her business cards
- Now, anyone who wishes to communicate privately with Alice
can look up her public key and proceed as above.
- Note that multiple senders can communicate multiple times
with Alice using the same public key pk for
encrypting all their communication.
- Note that pk is inherently public—
and can thus be learned easily by an attacker—in either of the above scenarios.

In the first case

an adversary eavesdropping on the communication between Alice and Bob obtains pk directly
in the second case
an adversary could just as well look up Alice's public key on its own.

We see that the security of public-key cannot rely on secrecy of pk
and must rely on secrecy of sk

It is therefore crucial
that Alice does not reveal her private key
to anyone
including the sender Bob.

Comparison to Private-Key Encryption

- Perhaps the most obvious difference
b/w private and public key encryption is that
the former assumes **complete secrecy**
of all cryptographic keys
whereas the latter requires
secrecy for only the **private key sk** .
- Although this may seem like a minor distinction
the ramifications are huge:
in the private-key setting
the communicating parties must somehow be able to
share the secret key without allowing any third party to learn it
while in the public-key setting
the public key can be sent from
one party to the other over a public channel
without compromising security.
- For parties shouting across a room
(or more realistically, communicating over a public network
like a phone line or the Internet)
public-key encryption is the only option.
- Another important distinction is that
private private key encryption schemes
use the **same key for both encryption and decryption**
while public-key encryption schemes
use different keys for each operation.
- I.e. public key encryption is inherently asymmetric.
- This asymmetry in the public key setting
means that the roles of **sender and receiver are**
not interchangeable
as they are in the private-key setting:
a single key-pair allows communication in one direction only
(bidirectional communication can be achieved in a number of ways

—the point is that a single invocation of a public-key encryption scheme forces a distinction b/w the receiver and the sender)

- In addition
 - a single instance of a public-key encryption scheme
 - enables multiple senders**
 - to communicate privately with a single receiver
 - in contrast to the private-key case
 - where a secret key shared b/w two parties
 - enables private communication among only those two parties.
- **Summarising** and elaborating the preceding discussion
 - we see that public key encryption has the following advantages relative to private-key encryption
 - Public key encryption addresses (to some extent) the key-distribution problem since communicating parties do not need to secretly share a key in advance of their communication. Two parties can communicate secretly even if all communication among them is monitored.
 - When a single receiver is communicating with N senders (e.g. an online merchant processing credit card orders from multiple purchasers) it is much more convenient for the receiver to store a single private key sk rather than to share store and manage N different secret keys (i.e. one for each sender).

In fact, when using public-key encryption the number of identities of potential senders need not be known at the time of key generation.

This allows enormous flexibility in "open systems".

The fact that **public-key encryption** schemes allow anyone act as a sender can be a **drawback** when a receiver only wants to **receive** messages **from one specific individual**
In which case, an authenticated (private-key) encryption scheme would be a better choice than public key encryption.

The **main disadvantage** of public-key encryption is that it is roughly **2 to 3 orders of magnitude slower** than private-key encryption (*This is an estimate, of course).

It can be a challenge to implement public-key encryption in severely resource-constrained devices

such as smartcards or RFID tags.

Even when a desktop computer is performing cryptographic operations carrying out thousands of such operations per second (as in the case of an online merchant processing credit card transactions) may be prohibitive.

Thus, when private key encryption is an option (i.e. if two parties can securely share a key in advance) then it typically should be used.

In fact as we will see in § 11.3 private-key encryption is used *in the public key setting* to improve efficiency for the (public-key) encryption of long messages.

A thorough understanding of private-key encryption is therefore crucial to appreciate how public-key encryption is implemented in practice.

Secure Distribution of Public Keys

- In our entire discussion thus far we have implicitly assumed that the adversary is passive.
i.e. the adversary only eavesdrops on communication between the sender and the receiver but does not actively interfere with the communication.
 - If the adversary has the ability to tamper with all communication b/w the honest parties share no keys in advance then privacy simply cannot be achieved.
 - For example if a receiver Alice sends her public key pk to Bob but the adversary replaces it with a key pk' of his own (for which it knows the matching private key sk') then even though Bob encrypts his message using pk' the adversary will easily be able to recover the message (using sk').
 - A similar attack works if an adversary is able to change the value of Alice's public key that is stored in some public directory

or if
the adversary can tamper with the public key
as it is transmitted from the public directory
to Bob.

- If Alice and Bob
don't share any information in advance
and are not willing to rely on some mutually trusted third party
there is nothing Alice or Bob can do
to prevent active attacks of this sort
or even to tell that such an attack is taking place

²In our "shouting-across-a-room" scenario, Alice and Bob can detect when an adversary interferes with the communication. But this is only because: (1) the adversary cannot prevent Alice's messages from reaching Bob, and (2) Alice and Bob "share" in advance information (e.g., the sound of their voices) that allows them to "authenticate" their communication.

- Importantly
our treatment of public-key encryption in this chapter
assume that senders are able to obtain a legitimate copy of the
receiver's public key.

(This will be implicit in the security definitions we provide).
i.e. we **assume secure key distribution**.

This assumption is made not because
active attacks of the type discussed above are of no concern
—in fact they represent a serious threat
that must be dealt with in any real-world system
that uses public-key encryption.

Rather
this assumption is made because there exist other
mechanisms for preventing active attacks
(e.g. Section 12.7) and it is therefore convenient (and useful)
to decouple the study of secure public-key encryption
from the study of secure public-key distribution.

11.2 Definitions

- We begin
by defining the syntax of public-key encryption
 - The definition is very similar to Definition 3.7 with
the exception that instead of working with just one key
we now have distinct encryption and decryption keys.

Definition 11.1

A *public-key encryption scheme* is a triple of PPT algorithms $(\text{Gen}, \text{Enc}, \text{Dec})$ such that

1. The *key generation algorithm* Gen
Input: the security parameter 1^n
Output: outputs a pair of keys (pk, sk) .

We refer to the first of these as the *public key* and the second as the *private key*.

Assume for convenience that pk and sk each has length at least n and that n can be determined from pk, sk

2. The *encryption algorithm* Enc
Input:
a public key pk and
a message m from some message space
(that may depend on pk).
Output
ciphertext c

We write $c \leftarrow \text{Enc}_{pk}(m)$
(looking ahead,
 Enc will need to be probabilistic to achieve
meaningful security).

3. The *deterministic decryption algorithm* Dec
Input:
private key sk
ciphertext c
Output:
a message m or a special symbol \perp
(denoting failure)

Write it as $m := \text{Dec}_{sk}(c)$.

It is required that,
except possibly with negl probability over (pk, sk)
output by $\text{Gen}(1^n)$
we have
$$\text{Dec}_{sk}(\text{Enc}_{pk}(m)) = m$$

for any (legal) message m .

< should it have
mentioned the
probability over Enc
as well?
Because Enc is also
PPT

The important difference from the private-key setting is that
the key generation algorithm Gen now outputs
two keys instead of one.

The public key pk is used for encryption

while the private key sk is used for decryption.

Reiterating our earlier discussion

pk is assumed to be widely distributed so that anyone can
encrypt messages for the party who generated this key
but sk must be kept private by the receiver in order for security to hold.

We allow for a negligible probability of decryption error
and
indeed

some of the schemes we present will have a negligible error prob
(e.g. if a prime needs to be chosen
but with negligible prob a composite
is obtained instead)

Despite this
we will generally ignore the issue from here on.

For practical usage of public-key encryption
we will want the message space to be $\{0,1\}^n$ or $\{0,1\}^*$
(and in particular, to be independent of the public key)

Although
we will sometimes describe encryption schemes using
some message space \mathcal{M} that does not contain all bit strings of some fixed length
(and that may also depend on the public key)
we will in such cases also specify
how to encode bit strings as
elements of \mathcal{M}

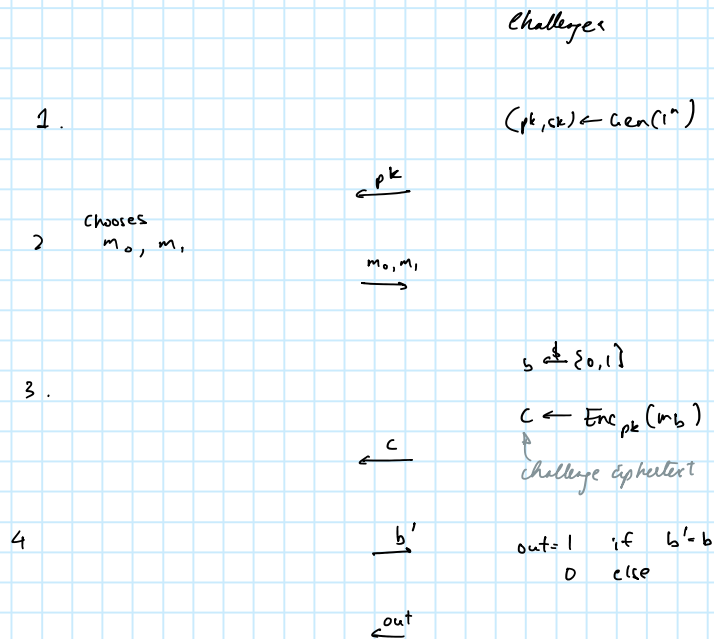
This encoding must be both
efficiently computable and efficiently reversible
so the receiver can recover the bit string that was encrypted.

11.2.1 Security against Chosen-Plaintext Attacks

- We initiate our treatment of security
by introducing the "natural" counterpart of Definition 3.8 in the
public-key setting.
 - Since extensive motivation for this definition
(as well as others) has been given in Ch 3
the discussion here will focus primarily on the differences
b/w the private key and public key settings.

- Given a public key encryption scheme $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$ and an adversary \mathcal{A} consider the following experiment:

The eavesdropping indistinguishability experiment $\text{PubK}_{\mathcal{A}, \Pi}^{\text{eav}}(n)$



Definition 11.2

A public-key encryption $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$ has *indistinguishable encryptions* in the presence of an eavesdropper if for all PPT adversaries \mathcal{A} there is a negligible function negl such that

$$\Pr[\text{PubK}_{\mathcal{A}, \Pi}^{\text{eav}}(n) = 1] \leq \frac{1}{2} + \text{negl}(n)$$

Story:

- The main difference b/w the above and Defn 3.8 is\ here \mathcal{A} is given the public key pk
 - Furthermore we allow \mathcal{A} to choose its messages m_0 and m_1 based on this public key.
 - This is essential when defining security of public-key encryption since as discussed previously

we assume that the adversary knows the public key of the recipient.

- The seemingly "minor" modification of giving the adversary pk has a tremendous impact:
 - it effectively gives \mathcal{A} access to an encryption oracle *for free*
 - (The concept of an encryption oracle is explained in Section 3.4.2)
 - The upshot is that Definition 11.2 is **equivalent to CPA-security**

We thus have

Proposition 11.3

If a public-key encryption scheme has indistinguishable encryptions in the presence of an eavesdropper it is CPA secure.

- This is in contrast to the private-key setting where there exist schemes that have indistinguishable encryptions in the presence of eavesdropper but are insecure under a CPA attack (see Prop 3.20).
 - Further differences from the private-key setting that follow almost immediately as consequences of the above are discussed next.

Impossibility of perfectly secret public-key encryption.

- Perfectly secret public-key encryption could be defined analogously to Definition 2.3 by conditioning on the entire view of an eavesdropper (i.e. including the public key).
 - Equivalently, it could be defined by extending Definition 11.2 to require that for *all* adversaries \mathcal{A} (not only efficient ones) it holds that

$$P_{\lambda} \left[\text{PubK}_{\mathcal{A}, \Gamma}^{\text{eav}}(n) = 1 \right] = \frac{1}{2}$$

- In contrast to the private-key setting however

perfectly secret public key encryption is *impossible*
regardless of how long the keys are
or
how long the message space is.

- In fact, an unbounded adversary given pk and a ciphertext c via $c \leftarrow \text{Enc}_{pk}(m)$ can determine m with probability 1.

A proof of this is left as Exercise 11.1.

Insecurity of deterministic public-key encryption.

- As noted in the context of private-key encryption no deterministic encryption scheme can be CPA secure.

Theorem 11.4

No deterministic public-key encryption scheme is CPA secure.

Story:

- Theorem 11.4 is so important it merits a bit more discussion.
 - The theorem is not an "artefact" of our security definition or an indication that our definition is too strong.
 - *Deterministic public-key encryption* schemes are vulnerable to practical attacks in realistic scenarios and should never be used.
 - The reason is that a deterministic scheme not only allows the adversary to determine when the same message is sent twice (as in the private-key setting) but also allows the adversary to recover the message with prob 1 (if the set of possible messages being encrypted is small!)
 - E.g. consider professor encrypting students' grades here an eavesdropper knows that each student's grade must be one of $\{A, B, C, \dots F\}$.

If the professor uses a deterministic public-key encryption scheme an eavesdropper can quickly determine any student's actual grade

- Although the above theorem seems deceptively simple for a long time *many real world systems were designed using deterministic public-key encryption.*

When public-key encryption was introduced
it is fair to say that the importance of prob. encryption
was not yet fully realised

The seminal work of Goldwasser and Micali
in which
(something equivalent to) Definition 11.2
was proposed Theorem 11.4 was stated
marked a turning point in the field of cryptography.


- The importance of pinning down one's intuition
in a formal definition and
looking at things the right way
for the first time
—even if seemingly simple in retrospect—
should not be underestimated.

11.2.2 Multiple Encryptions

- As in Ch 3
it is important to understand the effect of using
the same key (in this case, the same public key)
for encrypting multiple messages.
- We could formulate security in such a setting
by having an adversary output two lists of plaintexts
as in Definition 3.19.
- For the reasons discussed in Section 3.4.2
 - however
we choose instead to use a definition in which
the attacker is given access to a "left-or-right" oracle
 $LR_{pk,b}$ that
on input a pair of equal length messages m_0, m_1
computes the ciphertext $c \leftarrow \text{Enc}_{pk}(m_b)$ and
returns c .

Formally, consider the following experiment defined for a public-key encryption scheme
 $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$ and adversary \mathcal{A} :

The LR-oracle experiment $\text{PubK}_{\mathcal{A}, \Pi}^{LR-cpa}(n)$:



$$\begin{array}{c} b' \\ \hline \rightarrow \\ \leftarrow \text{out} \end{array} \quad \text{out} = \begin{cases} 1 & \text{if } b' = b \\ 0 & \text{else} \end{cases}$$

Definition 11.5

A public-key encryption scheme $\Pi = (\text{Gen}, \text{Enc}, \text{Dec})$
 has indistinguishable multiple encryptions if
 for all PPT adversaries \mathcal{A}
 there is a negligible function negl such that

$$\Pr[\text{PubK}_{\mathcal{A}, \Pi}^{\text{LR-epo}}(n) = 1] \leq \frac{1}{2} + \text{negl}(n)$$